

Lab 7: Programming with MPI

If you haven't completed Lab 1, review its handout and set up the environment for running MPI.

1 Simple Send-Receive

1. Read and understand the provided program `simple.c`. It implements a pair of send-receive actions between two processes. The sender sends an integer to the receiver; the receiver decreases the value by one and sends it back. Even though the message-passing happens only between two processes, the program can run with any number of processes. It can also take a command-line argument, an integer to be used as the message value. Compile and run it. Here are some examples:

```
linux> mpicc -o simple simple.c      // compile
linux> mpirun -n 2 simple            // run with 2 procs
linux> mpirun -n 4 simple            // run with 4 procs
linux> mpirun -n 4 simple 100        // run with 4 procs, msg=100
```

2. Write a program `ring.c` based on the above program. Instead of send-receive actions between two processes, the new program will involve *all* active processes. In the program, process 0 (*i.e.* process with `rank==0`) sends an integer to process 1; upon receiving the integer, process 1 decreases its value by 1, and sends the new number to process 2; process 2 does the same, and sends a new number to process 3; and so on. The last process in the active set sends its modified number back to process 0. Like in `simple.c`, each process should make the sending and receiving actions visible by printing out a message showing its rank, its host name, and the involved integer's value. Note that the total number of active processes is not controlled by the MPI program itself. Compile your program with `mpicc`, and test it with multiple combinations of runtime parameters.

2 Collective Communication

1. The file `vecsum.c` contains the vector sum program shown in class. It provides examples of various collective routines. Read and understand the program, then compile and run it.
2. The file `scan.c` contains a simple example of the prefix scan routine with the sum operation. Each process supplies its rank number as the source data and receives the scan result in a variable, `psum`. Compile and run this program, you should see each process prints out its `psum` value. For process i , the value is $\sum_0^i k$, *i.e.* the partial sum of all the rank numbers up to i .

Now, add a new gather routine, to collect all the `psum` values back to process 0. Have process 0 print out the array of received values from the gather routine, which should match those `psum` values.

3. The file `scatter.c` contain an example of the scatter routine. Process 0 scatters an array `cnt` of values to the other processes, value `cnt[i]` to process i . Compile and run this program.

Now, we want to add another scatter routine. This time we want to scatter a *different* number of items to each process. Follow the instructions below to change the program:

- (a) Add declarations:

```
int *data, *disp, rbuf[2 * size];
if (rank == 0) {
    data = (int *) malloc(sizeof(int) * 128); // source data
    disp = (int *) malloc(sizeof(int) * size); // displacement array
}
```

- (b) Have process 0 initialize the `data` array so that `data[k] = k` for all elements. This is the source array. We want to scatter a section of this array to each process; more specifically, we want to scatter `cnt[i]` number of `data` elements to process `i`.
- (c) Have process 0 initialize the `disp` array so that each entry holds a index value to the `data` array, indicating the start of a section. As an example, since `cnt[0] = 0; cnt[1] = 2; cnt[2] = 4;` (i.e. the first three section sizes are 0, 2, 4), we should have

```
disp[0] = 0; disp[1] = 0; disp[2] = 2; disp[3] = 6;
```

- (d) Add the following variable-sized scatter call to the program:

```
// scatter variable-sized data sections to all processes
MPI_Scatterv(data, cnt, disp, MPI_INT, rbuf, 2*size, MPI_INT, 0, MPI_COMM_WORLD);
```

Read MPI document to understand the parameters.

- (e) Add print statements after the call, so you'll see messages as follows:

```
P[0] got 0 items:
P[1] got 2 items: 0,1,
P[2] got 4 items: 2,3,4,5,
P[3] got 6 items: 6,7,8,9,10,11,
```

Compile and run the new program.

3 File I/O

1. Read and understand the program `file-in.c`. It opens a file; reads two integers from the file; and prints out their values. The input file name is provided as a command-line argument. Note that the input file is byte-encoded. To view its content, use the `od` command:

```
linux> od -i data.txt
```

Compile and run this program with different number of processes:

```
linux> mpirun -n 4 file-in data.txt
linux> mpirun -n 8 file-in data.txt
```

Modify the buffer size and have each process read in four integers.

2. Read and understand the program `file-out.c`. This program is similar to the previous one, except that it is for writing to files. Compile and run this program with any number of processes:

```
linux> mpirun -n 4 file-out output
```

What do you observe? How many output files are created? What are the contents?

Change this program so that there is only one output file, `output.all`; and have all processes write to this file. Compile and run. What do you see in this file? Can you explain?

3. Read and understand the program `file-view.c`. Pay attention to the `MPI_File_set_view()` line. Compile and run this program. Change the `offset` parameter and see the effect.
4. The program `sum-mpi.c` is copied from Lab 1. Modify this program to read input from a file. Specifically:

- (a) Comment out the routine `compute()`.
- (b) Change the parameter `N`'s value to 64.
- (c) After obtaining the values of `rank` and `size`, call `malloc()` to allocate an array to hold `N/size` integers.

- (d) Read `N/size` integers from file `data.txt` to this array.
 - (e) Compute the sum of the array's elements, stored the result in variable `psum`.
 - (f) Continue with the existing code to collect the partial sums into a global sum, and print it out.
- Compile and test your program.

4 Submission

As usual, write a short report, in plain text or pdf, summarize your work with this lab. Submit the report, `ring.c`, and the three modified programs, `scan.c`, `scatter.c`, and `sum-mpi.c`, through the “Lab7” submission folder on D2L (under the “Activities/Assignments” tab). You should submit your work before the week-end, *i.e.* Sunday 2/23.